

## Capstone 4 - "Fibre" (EUR/USD) Forex Analysis

### Background

The EUR/USD (Euro/US Dollar) currency pair is one of the most widely traded pairs in the foreign exchange (forex) market, representing the exchange rate between the Euro, the official currency of the Eurozone, and the US Dollar, the official currency of the United States. The popularity of this pair stems from the fact that it involves the two largest economies in the world, the Eurozone and the United States, making it a key indicator of global economic health and investor sentiment.

The dataset provided contains financial and trading-related data, including features such as **Trades**, **Gross Profit**, **Gross Loss**, **% Change**, **Profit Factor**, **Winners**, and **Net Profit**. This data represents key performance indicators (KPIs) that are typically used in evaluating the success of trading strategies and financial outcomes over a period of time. The dataset may include both summary statistics for trading strategies and profit metrics, which provide insights into profitability and risks associated with trades.

Given the financial nature of the dataset, it offers an opportunity to apply machine learning methods to derive valuable insights and predictive capabilities. Predictive modeling can aid in identifying trends, minimizing risks, and optimizing decision-making processes in financial trading strategies.

+ Code + Text

### Objective

The primary objective of this project is to apply various machine learning methods to analyze and predict financial performance metrics, such as Net Profit, based on features like Trades, Gross Profit, Gross Loss, % Change, and others. Specifically, the focus will be on the following:

1. **Prediction of Financial Outcomes:** Using supervised learning techniques such as **Regression Analysis** (e.g., Linear Regression, K-Nearest Neighbors, Neural Networks) to predict key financial metrics like Net Profit based on historical data.
2. **Exploration of Relationships Between Variables:** Utilize visualizations such as **pairplots** to explore the relationships between features and their impact on profit outcomes. This will help identify key drivers of profitability in the dataset.
3. **Improvement of Financial Forecasting:** Explore advanced machine learning models, including **Neural Networks**, to improve the accuracy of predictions, uncover hidden patterns, and provide actionable insights for optimizing trading strategies.

The overall goal is to enhance predictive accuracy, which can lead to better decision-making, risk management, and profitability in trading and financial strategy optimization.

### Origin of Data

Dataset is downloaded from [www.fxblue.com](http://www.fxblue.com) which I have collected the data between 8 Oct 2023 to 4 Aug 2024.

All the "Magic No." is an individual Expert Advisor (EA).

Dataset has the records of 200 rows with 40 columns.

```
1 # Importing basic libraries
2 import pandas as pd
3 import numpy as np
4 from sklearn.model_selection import train_test_split
5 from sklearn.linear_model import LinearRegression
6 from sklearn.tree import DecisionTreeRegressor
7 from sklearn.ensemble import RandomForestRegressor
8 from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
9 import matplotlib.pyplot as plt
10 import seaborn as sns
```

### Loading the Data

```
1 # Load the dataset
2 file_path = 'C:/KWM/NTUCLearning Hub/AI ML Developer Capstone Project/Forex.csv'
3 data = pd.read_csv(file_path)
4
5 # Display the first few rows of the dataset
6 data.head()
```

	Magic	Trades	Gross profit	Gross loss	Net profit	% change	Opening balance	Closing balance	Profit factor	Winners	...	Average winner length (hours)	Average loser length (hours)	Consec winners	Consec losses	Consec profit	Cor 1
0	1011000	388	279.79	-342.64	-62.85	-0.64	0	0	0.82	168	...	16.96	17.45	7	9	21.23	-1
1	1011001	229	432.87	-376.76	56.11	0.58	0	0	1.15	111	...	34.37	28.34	7	12	27.25	-3
2	1011002	242	426.83	-433.30	-6.47	-0.07	0	0	0.99	102	...	35.76	25.30	5	11	21.05	-3
3	1011003	321	361.47	-441.40	-79.93	-0.82	0	0	0.82	146	...	23.06	21.62	6	8	19.52	-2
4	1011004	266	310.55	-404.63	-94.08	-0.95	0	0	0.77	118	...	24.35	28.78	9	9	23.14	-4

## EDA and Data Preparation

```
1 data.info()
```

```
→ <class 'pandas.core.frame.DataFrame'>
RangeIndex: 200 entries, 0 to 199
Data columns (total 40 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   Magic            200 non-null    int64  
 1   Trades           200 non-null    int64  
 2   Gross profit    200 non-null    float64 
 3   Gross loss      200 non-null    float64 
 4   Net profit      200 non-null    float64 
 5   % change         200 non-null    float64 
 6   Opening balance 200 non-null    int64  
 7   Closing balance 200 non-null    int64  
 8   Profit factor   200 non-null    float64 
 9   Winners          200 non-null    int64  
 10  Winner%          200 non-null    int64  
 11  Losers           200 non-null    int64  
 12  Loser%           200 non-null    int64  
 13  Avg win          200 non-null    float64 
 14  Avg loss         200 non-null    float64 
 15  Avg trade        200 non-null    float64 
 16  Best trade       200 non-null    float64 
 17  Worst trade      200 non-null    float64 
 18  Longest (hours)  200 non-null    float64 
 19  Shortest (hours) 200 non-null    float64 
 20  Average length (hours) 200 non-null    float64 
 21  Total win pips   200 non-null    float64 
 22  Total loss pips  200 non-null    float64 
 23  Net pips          200 non-null    float64 
 24  Pip profit factor 200 non-null    float64 
 25  Avg win pips     200 non-null    float64 
 26  Avg loss pips    200 non-null    float64 
 27  Avg pips per trade 200 non-null    float64 
 28  Lots traded       200 non-null    float64 
 29  Avg lots          200 non-null    float64 
 30  Average winner length (hours) 200 non-null    float64 
 31  Average loser length (hours) 200 non-null    float64 
 32  Consec winners    200 non-null    int64  
 33  Consec losses     200 non-null    int64  
 34  Consec profit     200 non-null    float64 
 35  Consec loss       200 non-null    float64 
 36  Consec win pips   200 non-null    float64 
 37  Consec loss pips  200 non-null    float64 
 38  Valley (cash)     200 non-null    float64 
 39  Valley (pips)     200 non-null    float64 

dtypes: float64(30), int64(10)
memory usage: 62.6 KB
```

```
1 data.describe()
```

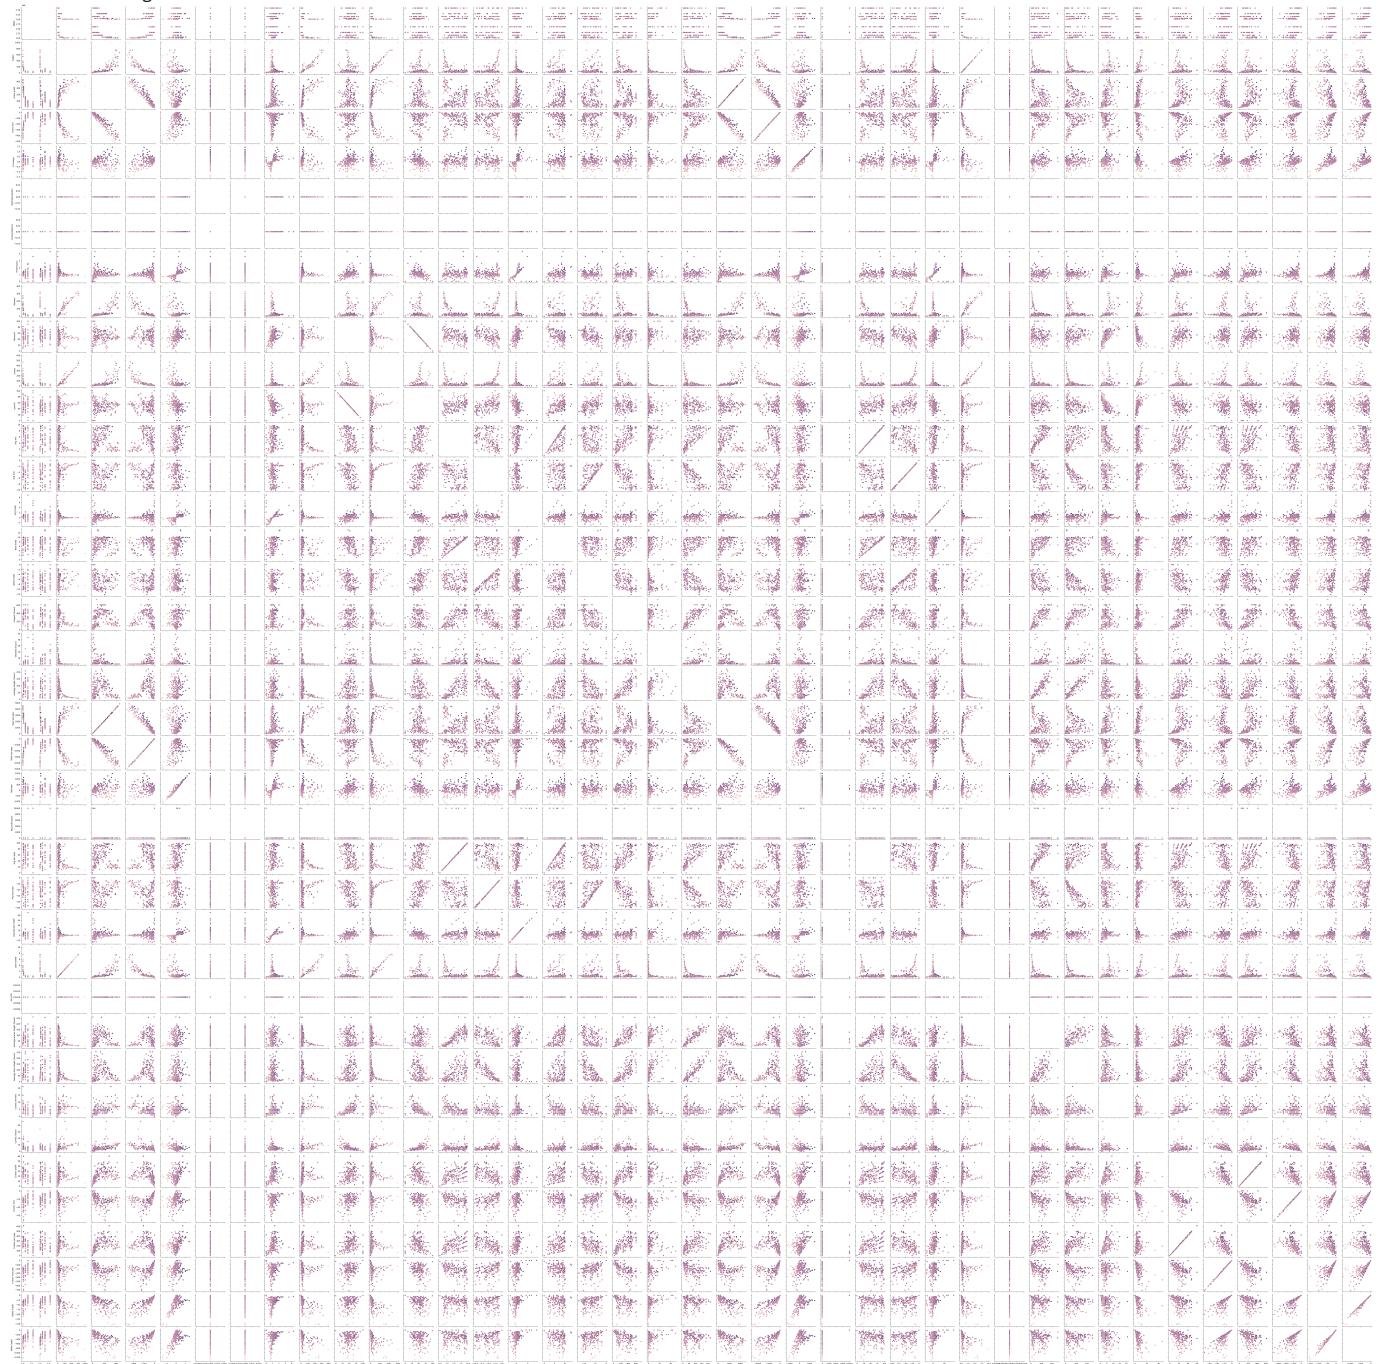


	Magic	Trades	Gross profit	Gross loss	Net profit	% change	Opening balance	Closing balance	Profit factor	Winners	...	Average winner length (hours)
<b>count</b>	2.000000e+02	200.000000	200.000000	200.000000	200.000000	200.000000	200.0	200.0	200.000000	200.000000	...	200.000000
<b>mean</b>	1.723504e+06	105.930000	152.990950	-154.669500	-1.678550	-0.016650	0.0	0.0	1.061200	50.250000	...	76.581150
<b>std</b>	5.451810e+05	161.929338	122.821435	135.082968	44.673946	0.452783	0.0	0.0	0.537034	73.684629	...	54.656456
<b>min</b>	1.011000e+06	1.000000	0.000000	-518.140000	-150.200000	-1.530000	0.0	0.0	0.000000	0.000000	...	0.000000
<b>25%</b>	1.151752e+06	23.000000	44.527500	-220.530000	-23.112500	-0.235000	0.0	0.0	0.790000	10.000000	...	26.385000
<b>50%</b>	1.806504e+06	41.500000	114.550000	-103.520000	0.140000	0.000000	0.0	0.0	0.990000	21.000000	...	72.005000
<b>75%</b>	2.151257e+06	88.000000	235.732500	-48.350000	22.445000	0.225000	0.0	0.0	1.242500	37.250000	...	116.970000
<b>max</b>	2.602009e+06	990.000000	460.400000	0.000000	138.800000	1.420000	0.0	0.0	4.290000	395.000000	...	259.700000

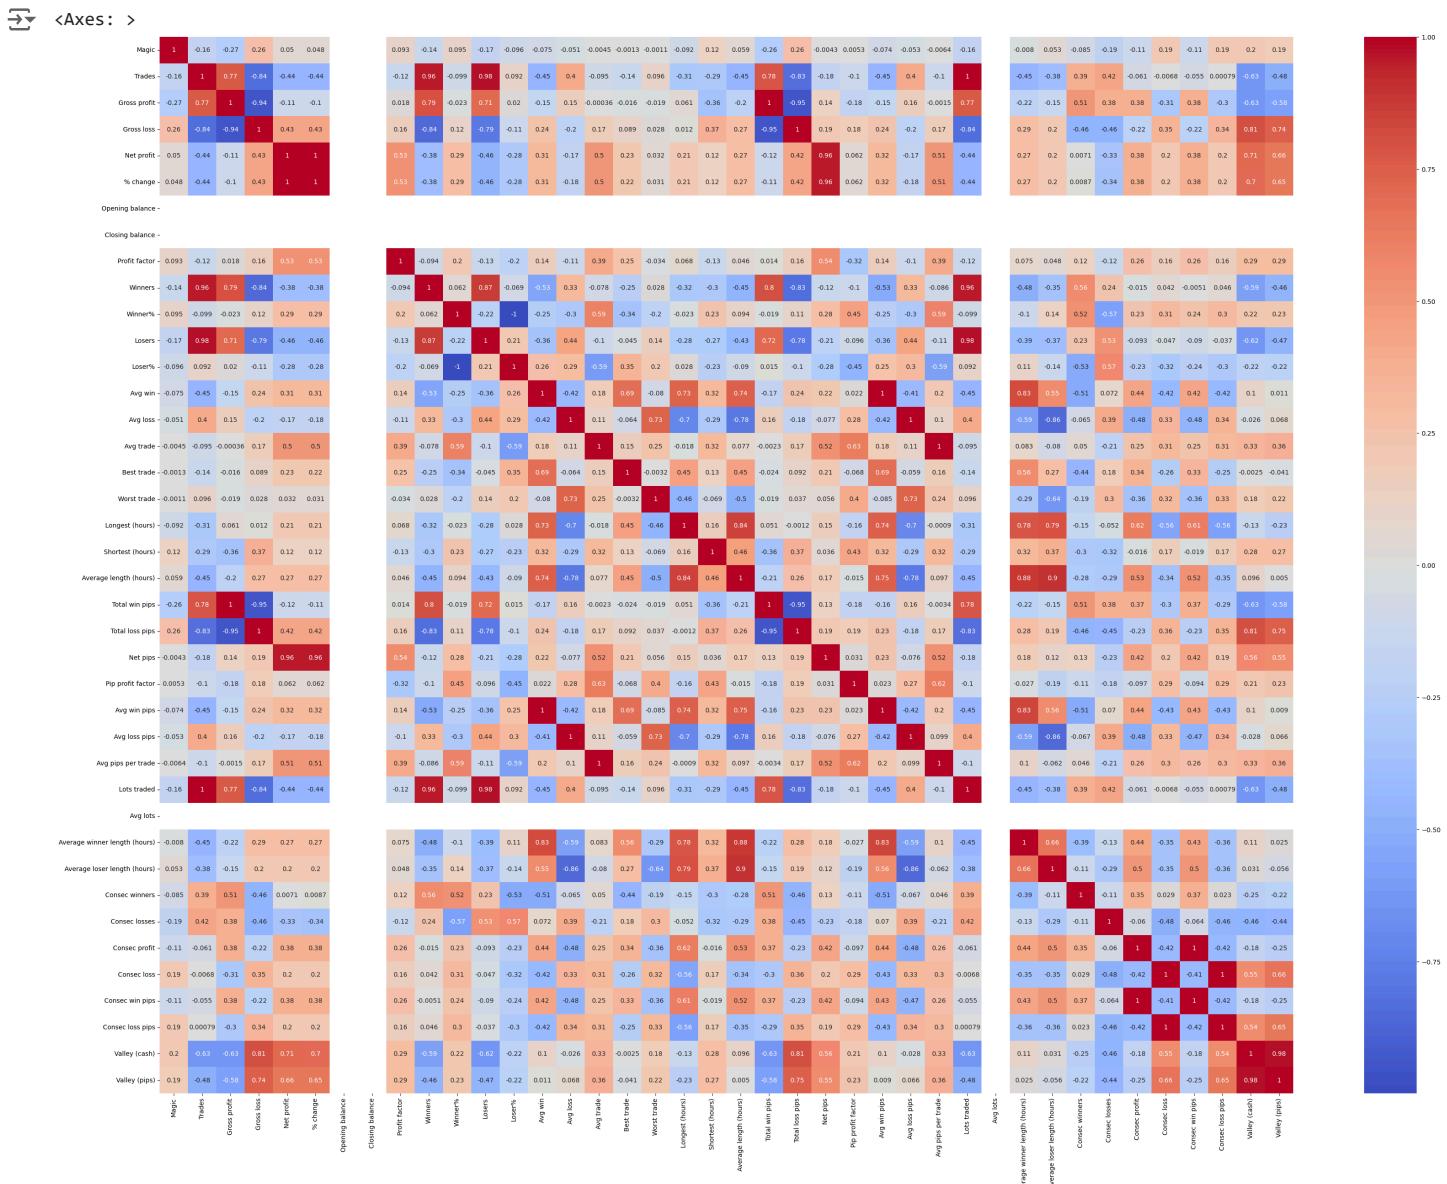
8 rows × 40 columns

```
1 sns.pairplot(data,hue='Net profit')
```

 <seaborn.axisgrid.PairGrid at 0x141f22d3910>



```
1 corrdf = data.corr()  
2 plt.figure(figsize=(40,30))  
3 sns.heatmap(corrdf, annot=True, cmap="coolwarm")
```



## Remove some of the columns

```
1 # Define the target variable (e.g., 'Net profit')
2 target = 'Net profit'
3
4 # Select features (excluding the target variable and any non-numeric columns)
5 features = data.drop(columns=['Average loser length (hours)', 'Average winner length (hours)', 'Avg pips per trade', 'Net pips', 'Magic',
6
7 # Convert all features to numeric if needed (useful if there are categorical variables)
8 features = features.apply(pd.to_numeric, errors='coerce')
9
10 # Drop any rows with missing values (optional, based on your data)
11 features = features.dropna()
```

```
1 features.head()
```

	Trades	Net profit	Profit factor	Winner%	Loser%	Avg trade	Best trade	Worst trade	Average length (hours)	Total win pips	Total loss pips	Pip profit factor	Avg win pips	Avg loss pips	Consec win pips	Consec loss pips
0	388	-62.85	0.82	43	56	-0.16	7.21	-7.59	17.10	2929.2	-3260.1	0.898500	17.4	-15.0	218.4	-147.9
1	229	56.11	1.15	48	52	0.25	4.10	-3.72	31.26	4434.9	-3675.3	1.206677	40.0	-31.1	279.6	-383.2
2	242	-6.47	0.99	42	58	-0.03	4.41	-3.80	29.71	4373.5	-4216.8	1.037161	42.9	-30.1	214.7	-330.0
3	321	-79.93	0.82	45	54	-0.25	4.89	-5.07	22.21	3745.4	-4273.8	0.876363	25.7	-24.6	198.7	-279.6

```
1 features.describe()
```

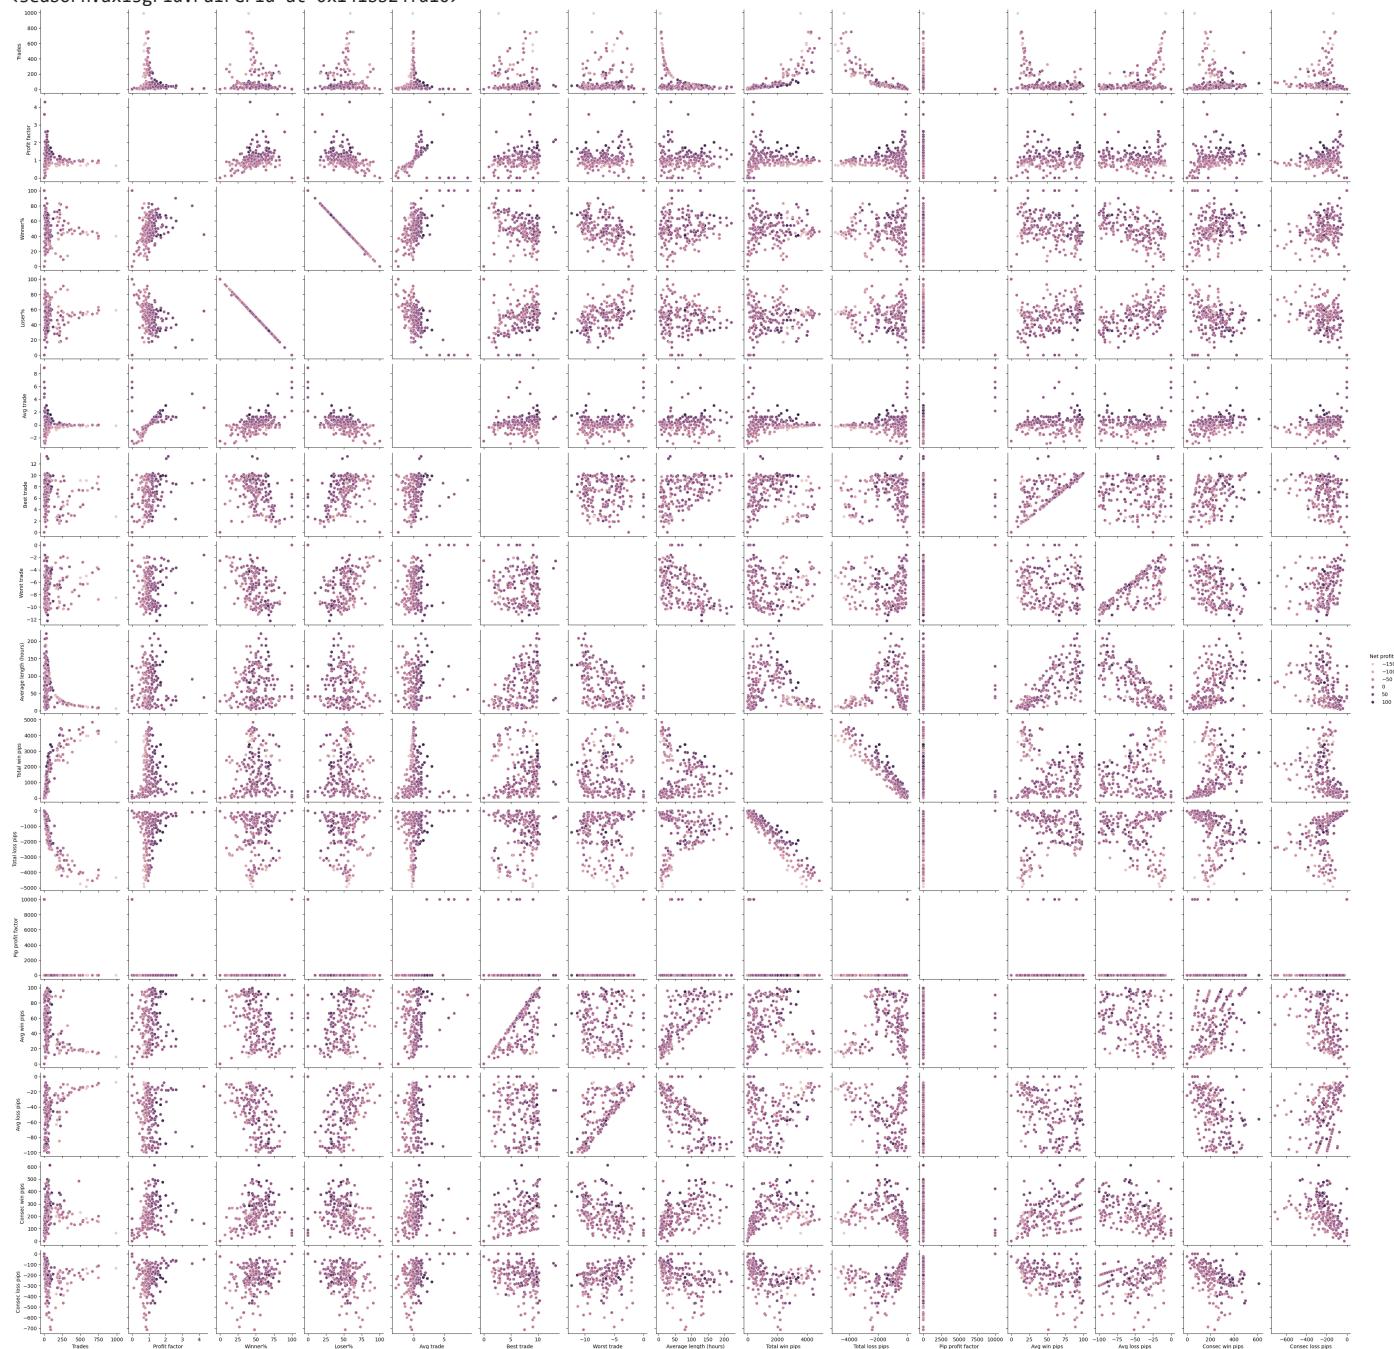
	Trades	Net profit	Profit factor	Winner%	Loser%	Avg trade	Best trade	Worst trade	Average length (hours)	Total win pips	Total loss pips
count	200.000000	200.000000	200.000000	200.000000	200.000000	200.000000	200.000000	200.000000	200.000000	200.000000	200.000000
mean	105.930000	-1.678550	1.061200	50.060000	49.865000	0.18385	6.934450	-6.815900	72.539800	1581.964000	-1499.750000
std	161.929338	44.673946	0.537034	17.701784	17.640086	1.37726	2.593167	2.748002	52.397489	1279.881718	1292.797626
min	1.000000	-150.200000	0.000000	0.000000	0.000000	-2.94000	0.000000	-12.250000	3.300000	0.000000	-4933.600000
25%	23.000000	-23.112500	0.790000	41.000000	37.750000	-0.40500	4.920000	-9.395000	26.960000	458.750000	-2156.225000
50%	41.500000	0.140000	0.990000	48.000000	52.000000	0.00000	7.405000	-6.785000	62.100000	1178.900000	-1027.450000
75%	88.000000	22.445000	1.242500	62.250000	59.000000	0.67000	9.115000	-4.700000	112.477500	2410.150000	-462.075000

```
1 features.info()
```

```
1 <class 'pandas.core.frame.DataFrame'>
RangeIndex: 200 entries, 0 to 199
Data columns (total 16 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   Trades          200 non-null    int64  
 1   Net profit      200 non-null    float64
 2   Profit factor   200 non-null    float64
 3   Winner%         200 non-null    int64  
 4   Loser%          200 non-null    int64  
 5   Avg trade       200 non-null    float64
 6   Best trade      200 non-null    float64
 7   Worst trade     200 non-null    float64
 8   Average length (hours) 200 non-null  float64
 9   Total win pips  200 non-null    float64
 10  Total loss pips 200 non-null    float64
 11  Pip profit factor 200 non-null  float64
 12  Avg win pips   200 non-null    float64
 13  Avg loss pips  200 non-null    float64
 14  Consec win pips 200 non-null    float64
 15  Consec loss pips 200 non-null    float64
dtypes: float64(13), int64(3)
memory usage: 25.1 KB
```

```
1 sns.pairplot(features,hue='Net profit')
```

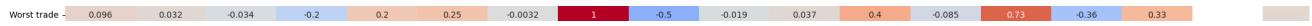
<seaborn.axisgrid.PairGrid at 0x141bb24fd10>



```
1 corrdf = features.corr()
2 plt.figure(figsize=(30,20))
3 sns.heatmap(corrdf, annot=True, cmap="coolwarm")
```



```
1 # Split the data into training and testing sets
2 X_train, X_test, y_train, y_test = train_test_split(features, data[target], test_size=0.2, random_state=42)
```



## Linear Regression Model



```
1 # Initialize the Linear Regression model
2 model = LinearRegression()
3
4 # Train the model on the training data
5 model.fit(X_train, y_train)
6
7 # Get the model's coefficients
8 print("Coefficients:", model.coef_)
9 print("Intercept:", model.intercept_)

→ Coefficients: [-1.94127432e-15  1.00000000e+00  2.57164994e-14 -4.97990519e-14
 -5.51045072e-14  1.81667403e-14 -7.15899562e-15 -1.45001997e-14]
```